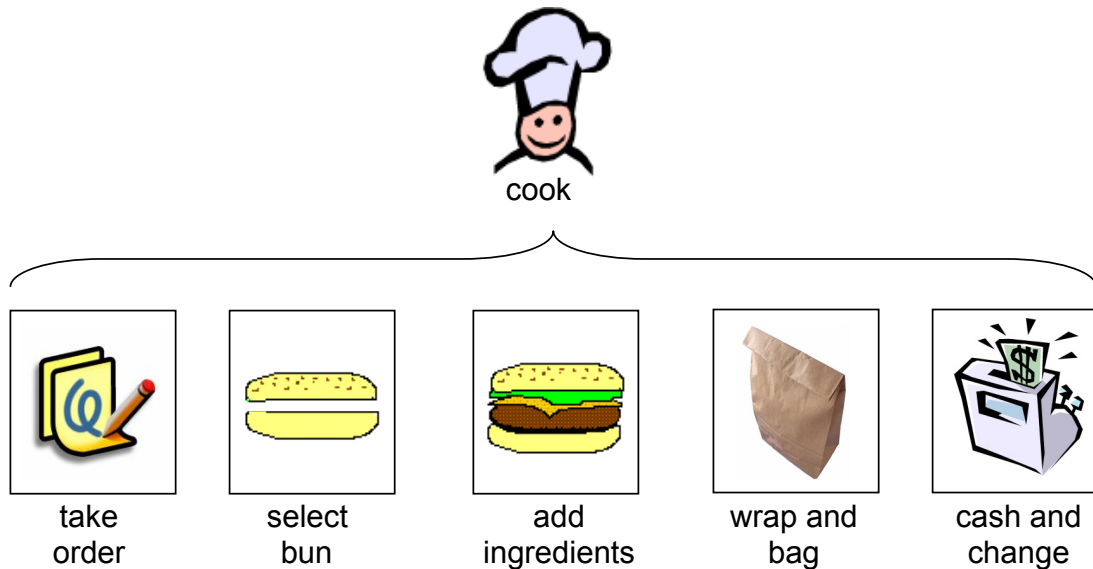


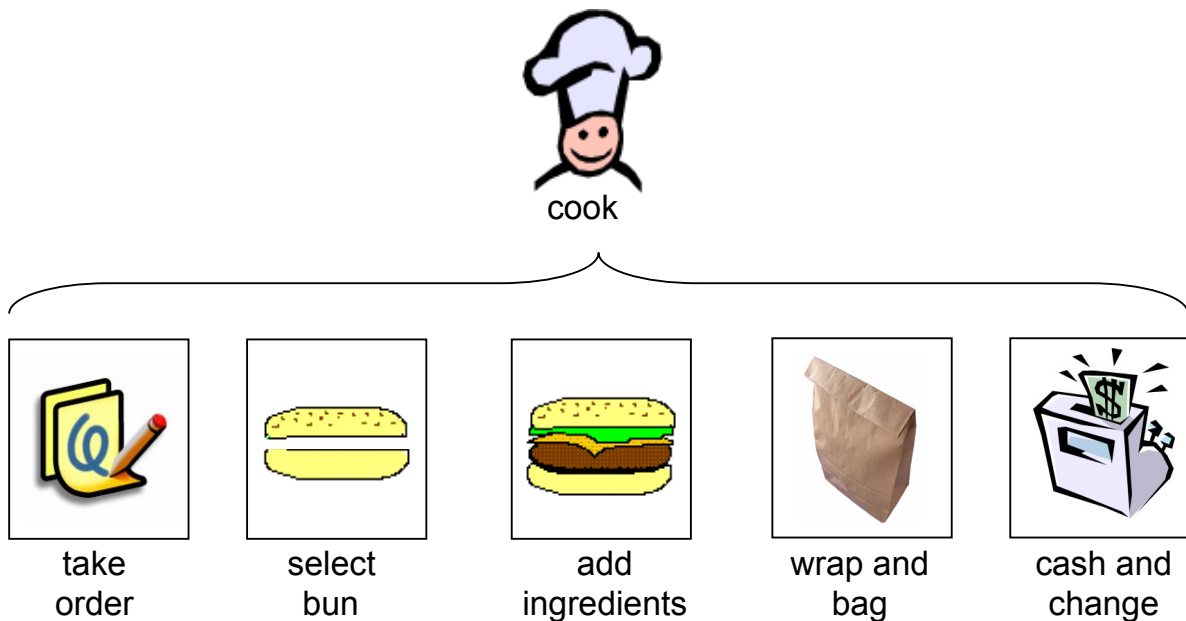
# ECE243

## CPU: Pipelining

### A Fast-Food Sandwich Shop



### With One Cook





customer1



customer1



customer1



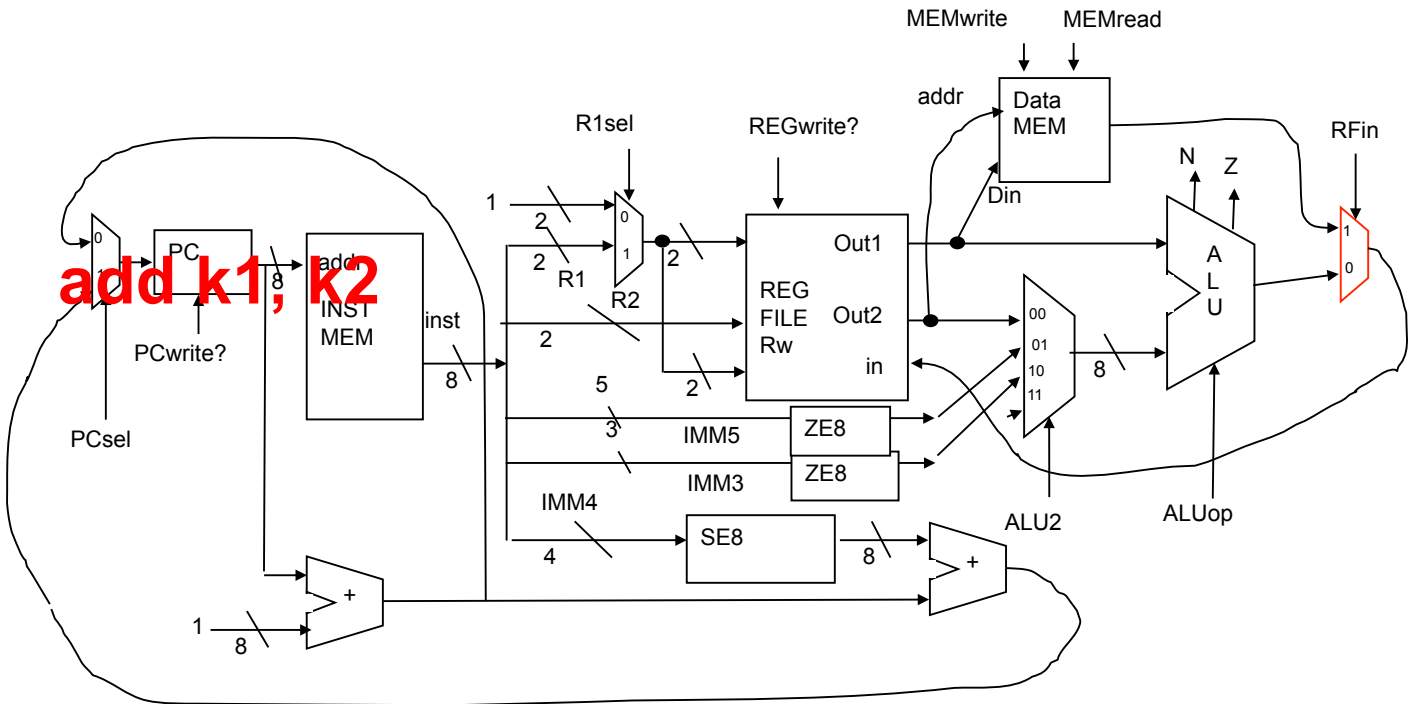
customer1



customer1

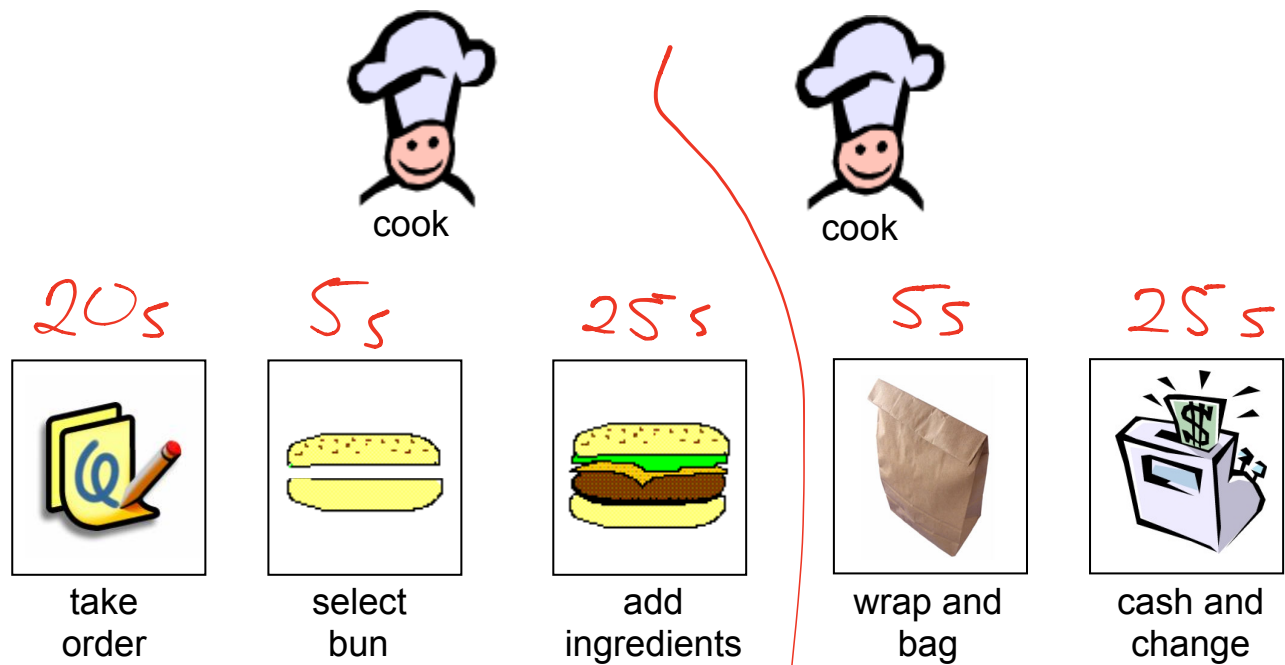
- one customer is serviced at a time

## Like the single-cycle CPU



- one instruction flows through at a time

# With Two Cooks?



Break up work so tasks are even

cook 1: 50s total

cook 2: 30s total

can service 1 new customer every 50s, 2 customers serviced @ a time

## Pipelining

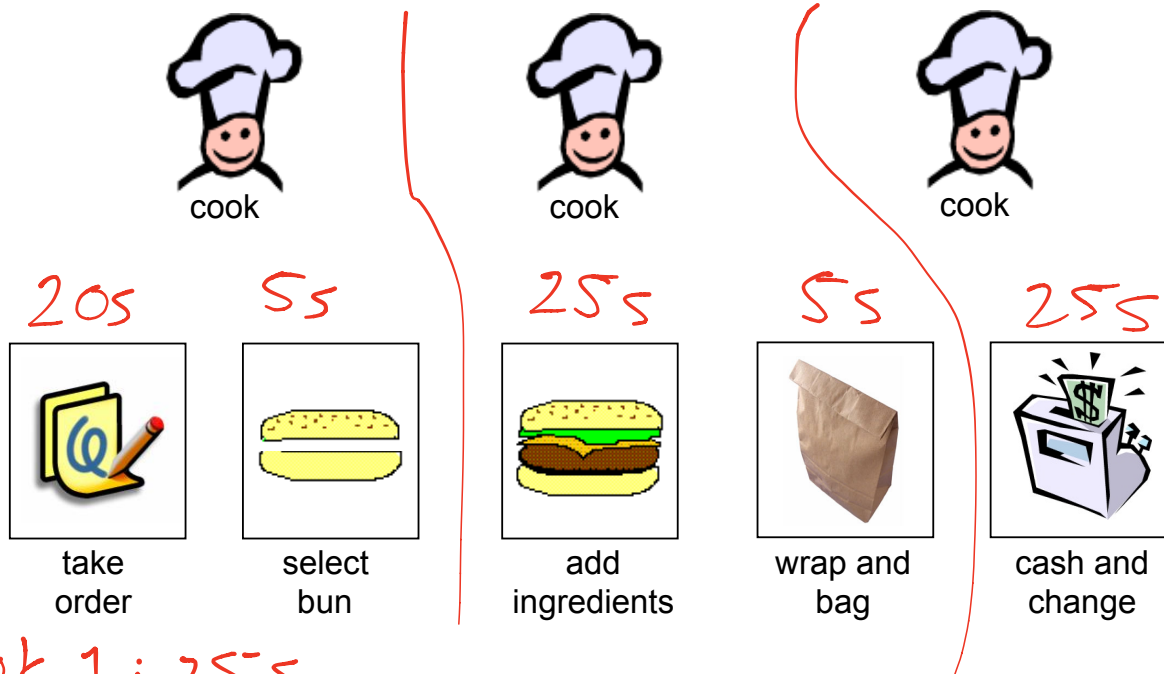
- Like an assembly line
- Doesn't change the interface or result
  - improves performance

## Pipelining a CPU (rough idea)



Insert pipe stage registers to store intermediate values  
modify control logic to support pipe stages

## With Three Cooks?



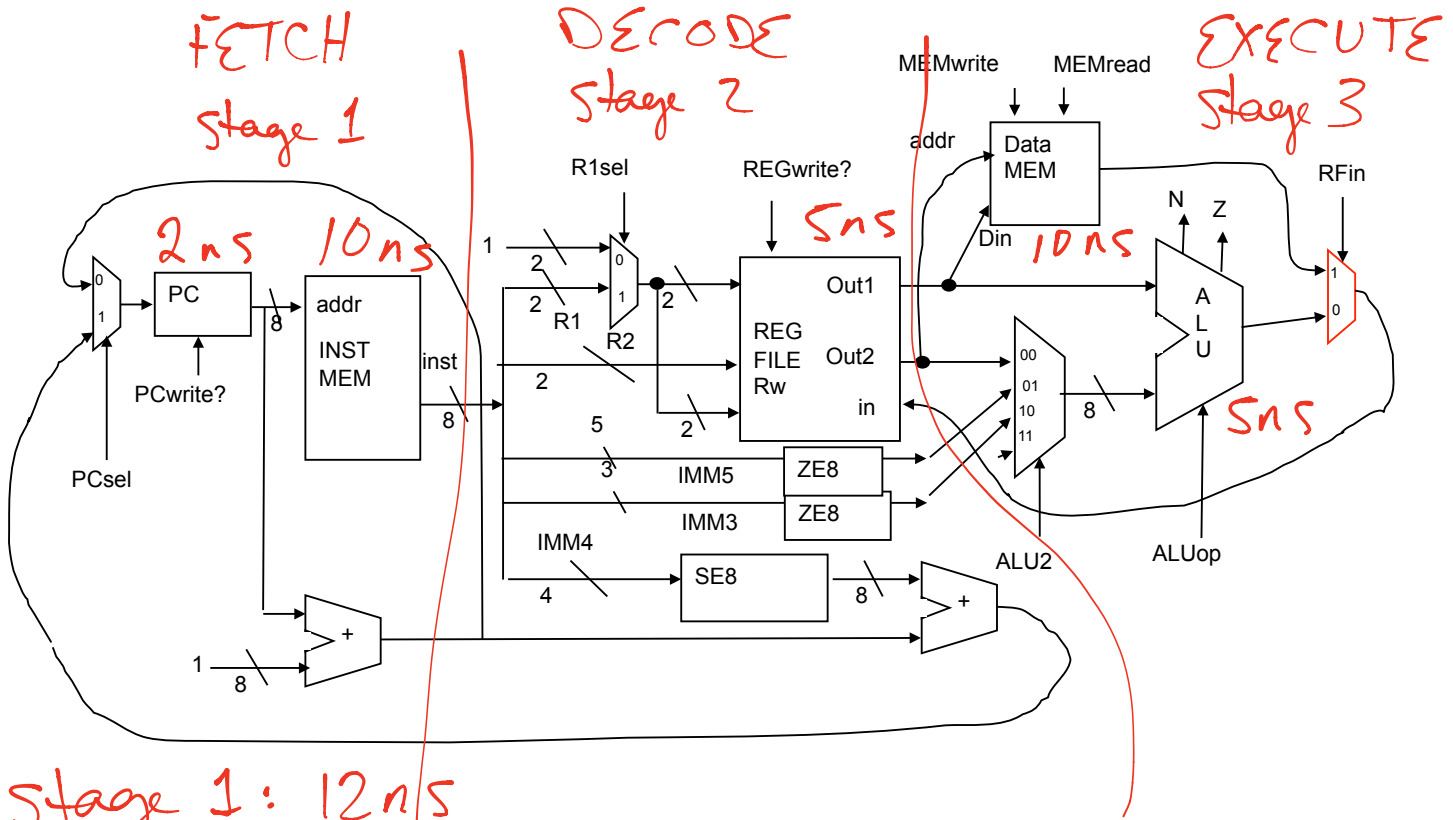
Cook 1: 25s

cook 2: 30s

cook 3: 25s

Can service new customer every 30s


## Pipelining a CPU (rough idea)



stage 1: 12ns  
 stage 2: 5ns  
 stage 3: 10ns


can start processing new instr every 12ns  
 3 instr at once

## Visualizing Pipelining



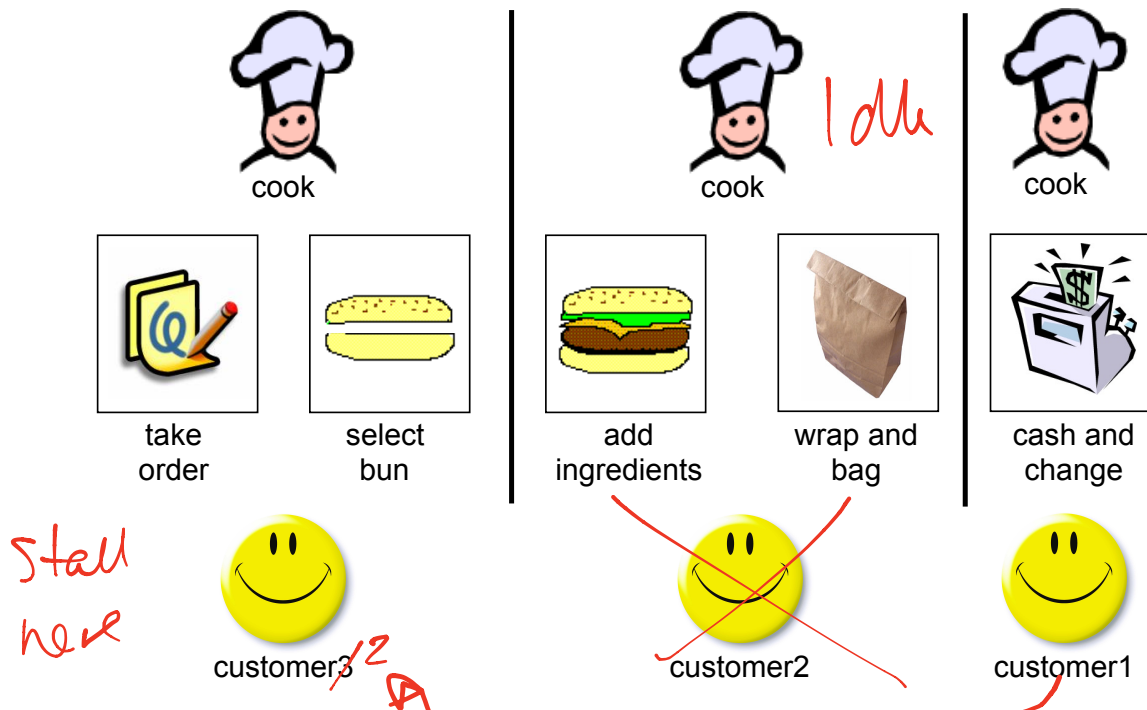
Cycle	Fetch	Decode	Execute
1	Instr 1		
2	Instr 2	Instr 1	
3	Instr 3	Instr 2	Instr 1
4	Instr 4	Instr 3	Instr 2

## Visualizing Pipelining (again)




Cycle	1	2	3	4	5
inst1	F	D	E		
inst2		F	D	E	
inst3			F	D	E
inst4				F	D

# Fast Food Hazards



What if: c1 and c2 are friends, c2 has no money, and c2 needs to know how much change c1 will get before ordering (to ensure c2 can afford his order)?

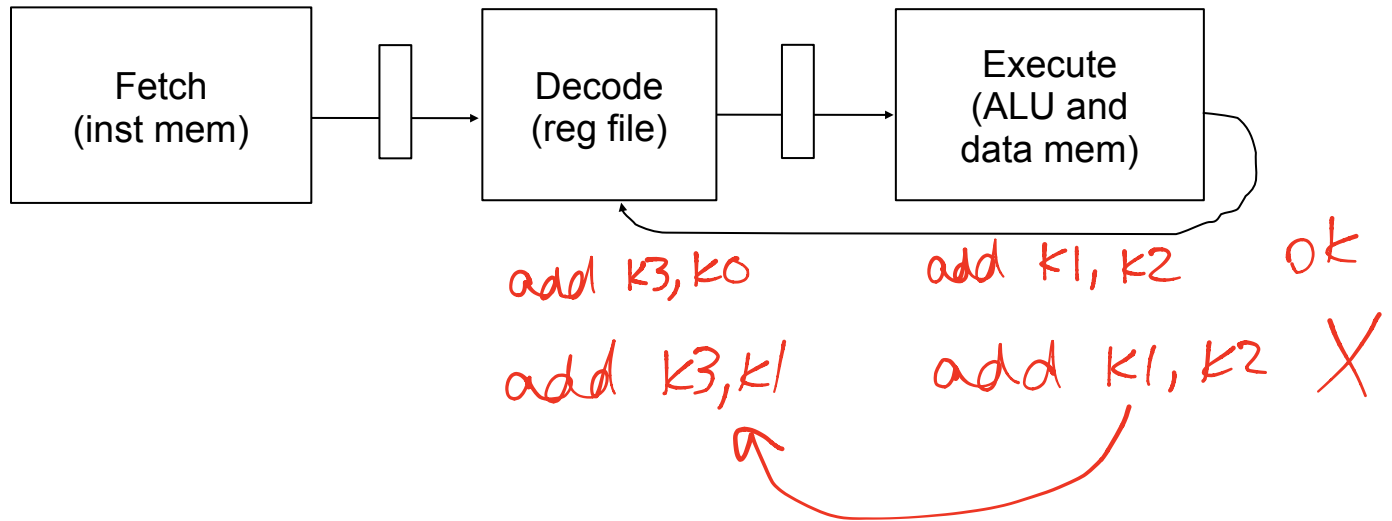
# Fast Food Hazards





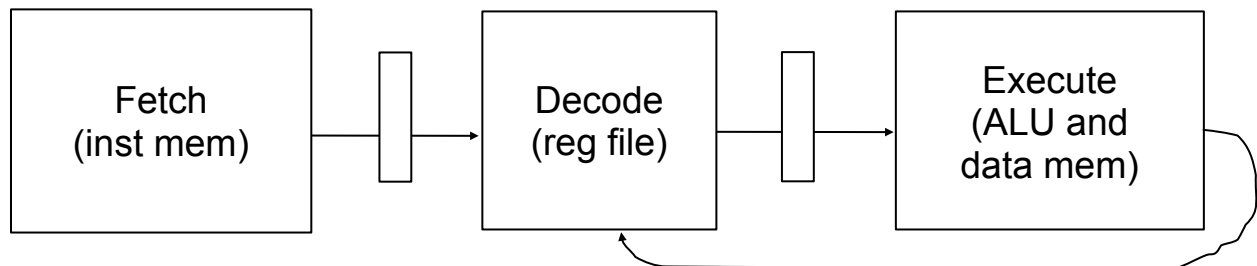
C2 will have to wait (aka stall) in 1st stage until C1 can tell him how much change he got

## CPU Hazards



- called a data hazard
- must be observed to ensure correct execution
- there are two solutions to data hazards

## Solution1: Stalling



Cycle	1	2	3	4	5
<code>add k1, k2</code>	F	D	E		
<code>add k3, k1</code>		F	F (stall)	D	E
<code>sub k0, k2</code>				F	D
<code>add k2, k2</code>					F

observe hazard by stalling  
insert bubble into pipeline

## How to insert bubbles

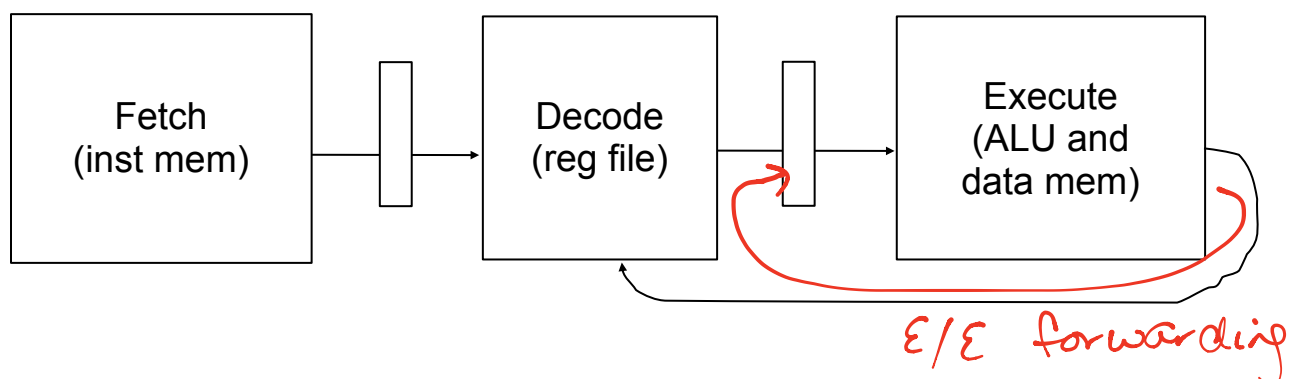
- option1: hardware stalls the pipeline
  - need extra logic to do so
  - happens ‘automatically’ for any code
- option2: compiler inserts “no-ops”
  - a no-op is an instruction that does nothing
  - ex: add r0,r0,r0 (NIOs)
  - compiler must do it right or wrong results!
- example: inserting a bubble with a no-op:

add k1, k2

nop

add k3, k1

## Solution2: Forwarding Lines



- add “forwarding” logic
  - to pass values directly between stages

Cycle	1	2	3	4	5
add k1,k2	F	D	E		
add k3,k1		F	D	E	
sub k0,k2			F	D	L
add k2,k2				F	D

## Control Hazards

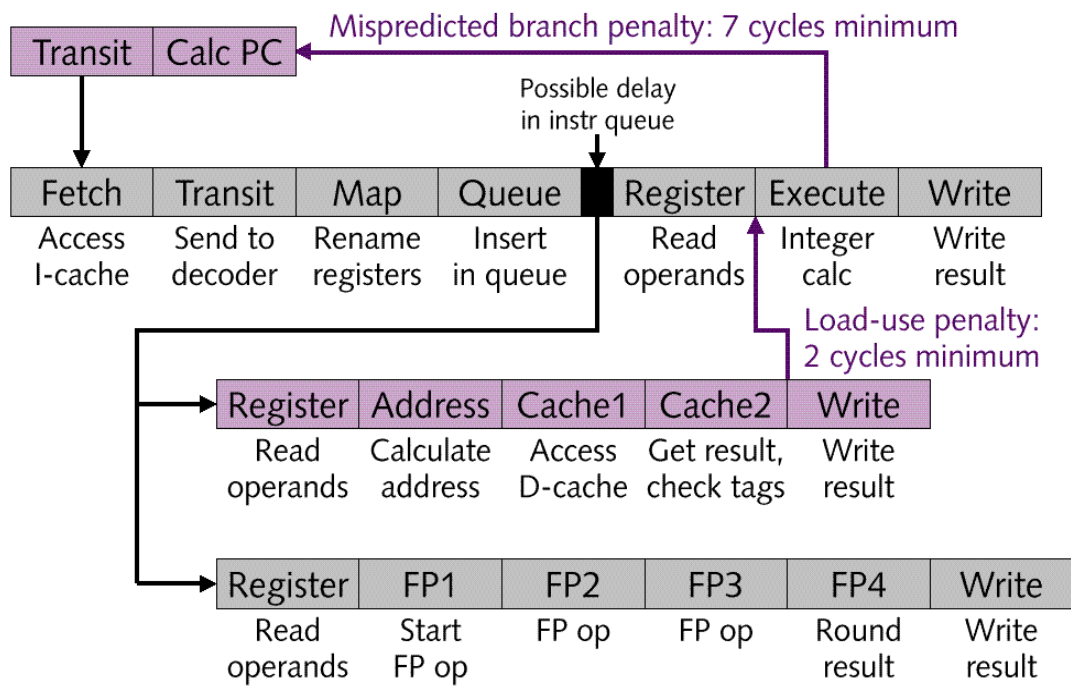
Cycle	1	2	3	4	5	6	7
add k1,k2	F	D	E		F	D	L
bnz -2		F	D	E		F	D
add k3,k1			F	D X			F
add k2,k2				F X			

- cpu predicts each branch is not taken
- Better: predict taken
  - why?---loops are common, usually taken
- More advanced: remember what each branch did last time

- “branch predictor”:
  - a table that remembers what each branch did the last time
  - uses this to make a prediction next time

# Some Real CPU Pipelines

## 21264 Pipeline (Alpha)



Microprocessor Report 10/28/96

## Pentium IV's Pipeline:

